PHIL 640: Supporting Notes for Four Arguments

Brian Weatherson

September 20, 2021

Relativism

The orthodox approach to semantics involves a commitment to something like the correspondance theory of truth. Declarative sentences are (typically) true or false, and when they are true, this is because they correspond to how things are. That last sentence needs a lot of spelling out. How can a sentence correspond to anything, apart from another sentence? Well, a sentence has a **content**. And that content can be true or false relative to various things. What things? Well, those are the matter of how things are in the platitude.

But how are things? It helps to work through an example. Let S be the sentence *Brian is sitting*. We will write $[\![S]\!]$ for the content of that sentence - these weird double square brackets have become the standard way for expressing the function from sentences to their contents. Now we want S to be true if, and only if, Brian is indeed sitting when S is uttered. That last bit, when S is uttered, is important. S doesn't say that I always sit, just that I'm now sitting. How should we incorporate that into our theory.

One approach is to build the time into the content of the utterance. On this view $[\![S]\!]_t$ is **Brian is sitting at t**, where t is the time the sentence was uttered. The subscript is to say that this is the content of S at t. On this view there is no such thing as *the* content of S, since it has different contents at different times of utterance. Then 'how things are' is the entire world, from Big Bang to Big Crunch. Call the world w, for simplicity, because we'll use it a lot. And we'll write $\vDash_w [\![S]\!]_t$ - which you can read as saying that the content of S is true at w, to say that this content is true. This approach is called **eternalism**, which is an annoyingly polysemous term in philosophy. If $[\![S]\!]_t$ is true, it is always true, hence the eternal. It's true even if I stand up, because $[\![S]\!]_t$ includes a reference to a time, and I stand at a later time. To introduce one more technical term, the time is part of the **context**; something that affects what the content of a given utterance of the sentence is.

Another approach is to build the time into how things are. On this approach the world does go from Big Bang to Big Crunch, but how things are doesn't just involve a world, it involves a time too. (If you like you can make this time metaphysically special - like on a spotlight view - but that's not part of the linguistic theory.) The content of $[\![S]\!]$ is just what it looks like: that Brian sits. But that very content is true at some times and false at others. We write this as $\vDash_{\langle w,t\rangle} [\![S]\!]$. This view is known as **temporalism**, because the things that are true or false are temporal entities, they change truth value over time. And to introduce another technical term, on this theory the **index** is the pair $\langle w, t\rangle$. The index is what the sentence is measured up against.

What is the difference between these views? There is a view that it's a purely notational difference, that we shouldn't care about which is right. But that's a minority view. Most philosophers think there is a big issue here. The problem is that there is less agreement about what exactly the big issue is. We want contents, the outputs of the $[\![]\!]$ function to play some philosophical role. The problem is that there are too many possible roles for them to play. Here are five possibilities.

1. The contents of what we believe.
2. The contents of what we say (if you think we have independent grasp on this).
3. The contents of 'that'-clauses, especially in indirect speech reports.
4. The things we agree or disagree with when we agree/disagree with an assertion.
5. The things we retract when we no longer endorse our former assertions.

Sometimes these five considerations push in the same direction. So consider the sentence *I was born in Australia*. That's true when uttered by me, and false (I think) when uttered by anyone else here. But when I say it, you don't disagree with me. And you don't report me by saying "Brian said that I was born in Australia". For reasons like this, everyone agrees that who the speaker is is relevant to the **content** of an utterance involving a first-person pronoun, not to the **index**. It's not that I say something that's true relative to me and false relative to you. It's that I say something that's true full stop, even though had you uttered the same words, you would have said something false. That's for the simple reason that you would have said something different.

But sometimes these considerations come apart. In the temporalism/eternalism debate, some considerations push towards putting time into the index, and some into the content. And so there's a debate here, but I'm not going to go deep into it.

Just one quick thing to note though. I've talked so far as if it is obvious that worlds go into indices. And I think it is obvious that's true. But not everyone agrees. Some philosophers have recently argued that worlds should go into contexts and hence into contents. So when I say "There is a famous detective that lives on Baker Street", that's false, but the same utterance would be true in Sherlock-Holmes-world. That's because, says orthodoxy, there is a proposition *That a famous detective lives on Baker Street* which is false in our world and true in Sherlock-Holmes-world. But some philosophers think that just like times go into contents (given eternalism), worlds should go into contents too. So when I say "There is a famous detective that lives on Baker Street", I express the proposition *That a famous detective lives on Baker Street in w*, where w is a name for our world. As I said, I think this is an absurd view, but it's worth noting it is out there. Still, most people think that indices are non-trivial, and they include at least worlds.

A huge amount of work in philosophy, especially in the 2000s, was on questions around the nature of context and index. In particular, these three questions took up a lot of people's time and attention. (The first by far the most, then the second, and then the third.)

1. Which utterances are such that their truth is sensitive to context? A huge debate, one of the biggest of the 90s and 2000s, was about whether knowledge ascriptions were among them, with **contextualists** (about knowledge) saying that utterances like "Brian knows that cars exist" could be true in regular contexts, but false in 'sceptical contexts'. But this was merely one manifestation of a discipline wide search for context-sensitivity.

2. If utterances are sensitive in this way, is this because the context affects their content, or because it affects the index they are evaluated against?
3. If it does affect the index, does it do so in a uniform way, or are there some sentences that interact differently with indices to others?

Don't worry about the last - I've got it there for completeness. It's mostly because the most important figure in this whole debate, John MacFarlane, disagrees with the way I've set up the debate so far. And it would be misleading to write his view out of the picture, since he is really the most important figure in it. MacFarlane's view is particularly motivated by issues about retraction, but I'm going to set those issues aside, and mostly set the precise views of MacFarlane aside as well.

For an example of what's at stake at least in questions 1 and 2, imagine that we have with us an ancient Spartan who very helpfully speaks English. And now let *S* be the sentence "Infanticide is wrong". We endorse that sentence, and the Spartan rejects it. Here are three possible philosophical takes on that dispute.

First, we could have an invariantist approach. One of us, presumably us, is right, and the other is wrong. This is the view that most moral realists take, though actually a lot of views that are not particularly realist could take it as well.

Second, we could have what I'd call a **contextualist** approach. When we say "Infanticide is wrong", we speak truly. But when the Spartan says "Infanticide is not wrong", he speaks truly as well. Why? Because our sentences have different contents. The content of our sentence is *That infanticide is wrong by contemporary standards*. And the content of the Spartan's sentence is *That infanticide is wrong by Spartan standards*. And those propositions are both true.

Third, we could have what I'd call a **relativist** approach. (Though note MacFarlane denies this is relativism; he calls this view 'non-indexical contextualism'.) Both we and the Spartan speak truly. But we also disagree, because we endorse the very proposition that the Spartan rejects. How is this possible? Well, the proposition we are disputing is only true relative to an ordered pair of a world and a moral standard. Relative to our world/standard, it is true; relative to the Spartan's, it is false. It's just like if I said "There is no famous detective living on Baker Street", that's true, even though it expresses the negation of a proposition that a person in Sherlock-Holmes-world can truly express. In a figurative sense, we and the Spartan live in different moral worlds, so our thoughts and talk can clash, without either being wrong.

This was a very long response to an argument Williams briefly makes against moral relativism. He thinks it is a problem for relativism that people in isolated moral communities won't know that there are other communities, and so won't know to relativise their moral talk. I think this is only a problem for moral *contextualism*, not moral relativism properly speaking. The real moral relativist, in the sense of the previous chapter, does not think that the moral utterances that people make are in any sense about moral communities, or standards, or anything of the sort. Rather, they think that the theory of truth for the (simple) propositions those utterances express, involves relativity to communities/standards. And they don't have to say that competent speakers need realise that, any more than competent speakers have to know anything about the physics of time in order to use tensed language.

The Meta-Ethical Trilemma

McDowell presents the following as an argument for non-cognitivism, but I find it more helpful to consider it as a trilemma (which is of course a trivial variation on McDowell's presentation.)

1. Moral attitudes provide reasons.
2. Beliefs alone do not provide reasons.
3. Moral attitudes are beliefs.

There are three obvious ways to respond to this, by denying each of the three lines. McDowell calls denials of 1 **descriptivism**. Certain kinds of reductivism will end up in this camp (if they don't end up as error theories.) He doesn't give a label for denials of 2, but it's the position he wants to take. And denials of 3 are **non-cognitivism**.

There are a number of other ways to get out of the trilemma without cleanly denying any of the three steps.

- As always in philosophy, it is possible to say that the argument equivocates. Here the most obvious equivocation is on 'reason'. Perhaps there is a hidden equivocation between *motivating reason* and *justifying reason*. But a better way out is to say that there is an equivocation between *decisive reason* and *supporting reason*. Foot, for instance, denies 1 if 'reason' means *decisive reason*, and denies 2 if 'reason' means *supporting reason*.
- The whole setup assumes that attitudes are the kind of thing that provide reasons, and without that assumption (which is a lot less popular now than it was 40 years ago) it isn't clear there is a problem here.
- You could deny more than 1 of the statements. The view in my book entails that 1 and 2 are both false (while the view is neutral on 3). But that's way outside the mainstream meta-ethical views.
- You could deny that there is an inconsistency here to be resolved. I *think* you can develop a moral error theory that endorses all 3.

Beliefs and Reasons

McDowell offers two reasons for endorsing 2. He isn't maximally clear about them, but here is my best attempt at spelling them out in more detail. The first argument starts from what I think the 'hydraulic' conception of reasons is meant to be.

1. For the rational person, the reasons that motivate their actions are the reasons that justify their actions.
2. Beliefs are passive states; all it takes to believe something is to have a library in your head.
3. Motivating reasons are not passive.
4. So beliefs are not motivating reasons (from 2, 3).
5. So, for the rational person, beliefs are not justifying reasons.

I think it's plausible to reject both 1 and 2. A rational person need only be sensitive to their reasons, not be motivated by them. Reasons can serve as guardrails, stopping the person when they are going to do the wrong thing, without motivating them when they are doing the right thing. And 2 will be false on a **functionalist** picture of beliefs. (Though some people think this is a reason to reject functionalism.)

The other argument I think is more interesting. This is the argument from the idea that it must be possible to tell that a reason is a reason from a 'sideways on' perspective. That is, it must be possible to tell that it is a reason from a neutral perspective. This view finds its modern expression in theories that demand that norms (moral, epistemic, whatever) are *guiding*. Here's the argument I read McDowell as giving.

1. If A is a reason for X, it is possible to tell from a neutral perspective that A is a reason for X.
2. From a neutral perspective, it is not possible to tell that a belief (alone) is a reason for any action.
3. So, beliefs (alone) are not reasons for action.

This is a really interesting argument. I agree with McDowell both that 1 is false, and that a lot of philosophical mistakes are caused by thinking that 1 is true. One thing that you get from people believing 1 is you get a theory of reasons which is deeply tied up with formal notions of coherence. So for something to be a reason it has to logically entail the thing being done, or you have to be able to compute that the thing has maximal value, or something like that. These approaches seem false to me twice over. For one thing they licence all sorts of terrible - but utterly coherent - views. For another they don't even get the promised neutrality, since questions of logic, decision theory, and game theory, and hence of coherence, are just as hard to tell from the neutral perspective as moral questions.

There is another more contemporary view that actually agrees with McDowell that beliefs can be reasons, but which is motivated by the same kind of 'neutral perspective' consideration. This is the view that only thin moral beliefs can motivate. So a belief *A is wrong* is a reason not to do A, but literally no other kind of belief can be reason-providing. Again, I think this is driven by a desire that reasons are detectable as reasons from a neutral perspective.

Non-Cognitivism and Separation

I thought the argument in section 5 of McDowell's paper was really fascinating, but I'm not sure it has always been fully appreciated in the subsequent literature. (Not that I've read 100% of it, or even close.) The argument is an attack on the view that attitudes involving thick concepts, like say 'honest', involve two separate parts. On the one hand, they involve a belief that the honest thing satisfies the descriptive part of 'honest'. On the other, they involve some kind of pro-attitude towards the thing in virtue of satisfying that descriptive component.

One way to argue against this view would be to simply say that introspection does not reveal that the attitude is bifurcated this way. But that's a lousy argument I think. We shouldn't trust introspection that much. Here's a slightly better argument.

1. If 'honest' has an evaluative and a non-evaluative part, then someone who doesn't share our values could understand the non-evaluative part.
2. If so, they could predict how we would apply the term, even without sharing the values.
3. But someone who doesn't share our values can't predict that.
4. So 'honest' doesn't have these two parts.

But as Blackburn pointed out, in this argument premise 2 is false. If 'honest' means 'good in virtue of H', someone could understand H and not predict our usage of it. Here's a much more interesting argument that turns on the fact that neither we nor outsiders can factorise 'honest'.

1. From the inside, the evaluative part of 'honest' feels non-arbitrary; it doesn't seem like we just like honest things the way we might like a flavour of ice-cream. Call this the non-arbitrariness constraint.
2. The non-arbitrariness constraint requires that we can identify the descriptive features in virtue of which we positively evaluate them, and do so independently of our own evaluation.
3. We can only identify the descriptive parts of our own thick concepts independently of our own evaluations if an outsider, who did not share or understand our values, could identify them.
4. An outsider could not make this identification - this is both intuitive, and arguably the result of the long Wittgensteinian discussion in section 3.

And now we have a problem. Line 1 says that our own attitudes involving thick concepts satisfy the non-arbitrariness constraint, but lines 2-4 together imply that they could not.

This is why I think it matters that we can't identify the honest, or kind, or cruel, behaviors independently of our values. Given non-cognitivism, the only way we can pick out the kind behaviours are as those that morally resemble, by our lights, the paradigms of kind behaviour. But it would feel hopelessly circular to say that we admire kind behaviour because things only fall under the concept 'kind' if we admire them. So there is a problem for non-cognitvists here.

If I were a non-cognitivist, I'd try to get out of the problem this way. All that non-arbitrariness requires is that on any given occasion of use, we can specify in a non-circular way why we admire the people we admire. And that's possible even if we can't specify exactly what kindness is. But is this enough for the non-cognitivist?